

# CS 161 Fundamentals of Artificial Intelligence

## Lecture 2

### Problem Solving and Uninformed Search

Quanquan Gu

Department of Computer Science  
UCLA

Jan 12, 2023

# Outline

- Problem-solving agents
- Problem types
- Problem formulation
- Example problems
- Basic search algorithms

# Problem solving as search problem

- ▶ Many AI problems can be formulated as search
- ▶ For example, "Farmer Crosses River Puzzle"

# Problem-solving Agents - Example: Romania

- On holiday in Romania; currently in Arad.

Flight leaves tomorrow from Bucharest

Formulate goal:

be in Bucharest

Formulate problem:

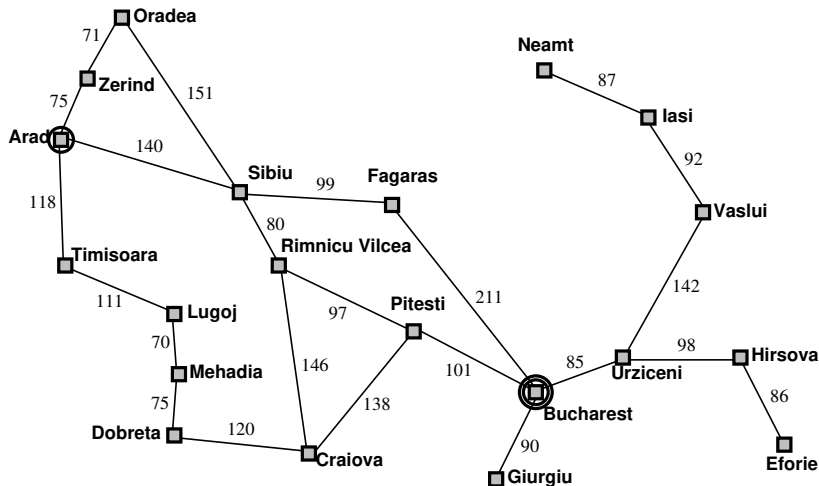
states: various cities

actions: drive between cities

Find solution:





sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest

## Example: Romania



# Example: vacuum world

Single-state, start in #5. Solution??



# Example: vacuum world


Single-state, start in #5. Solution??

[*Right*, *Suck*]

Conformant, start in {1, 2, 3, 4, 5, 6, 7, 8}

e.g., *Right* goes to {2, 4, 6, 8}.

Solution??



# Example: vacuum world

Single-state, start in #5. Solution??

[*Right, Suck*]

Conformant, start in {1, 2, 3, 4, 5, 6, 7, 8}

e.g., *Right* goes to {2, 4, 6, 8}.

Solution??


[*Right, Suck, Left, Suck*]

Contingency, start in #5

Murphy's Law: *Suck* can dirty a clean carpet

Local sensing: dirt, location only.

Solution??





# Example: vacuum world

Single-state, start in #5. Solution??

[*Right, Suck*]

Conformant, start in {1, 2, 3, 4, 5, 6, 7, 8}

e.g., *Right* goes to {2, 4, 6, 8}.

Solution??

[*Right, Suck, Left, Suck*]


Contingency, start in #5

Murphy's Law: *Suck* can dirty a clean carpet

Local sensing: dirt, location only.

Solution??

[*Right, if dirt then Suck*]



# Single-state problem formulation

A **problem** is defined by the following items:

**states** e.g., city names

**initial state** e.g., “at Arad”

**actions** e.g.,  $\langle Arad \rightarrow Zerind \rangle$

**successor function**  $S(x)$  = set of action–state pairs

e.g.,  $S(Arad) = \{\langle Arad \rightarrow Zerind, Zerind \rangle, \dots\}$

**goal test**, can be

**explicit**, e.g.,  $x = \text{“at Bucharest”}$

**implicit**, e.g.,  $NoDirt(x)$

**path cost** (additive)

e.g., sum of distances, number of actions executed, etc.

$c(x, a, y)$  is the **step/action cost**, assumed to be  $\geq 0$

A **solution** is a sequence of actions

leading from the initial state to a goal state

# Selecting a state space

Real world is absurdly complex

⇒ state space must be **abstracted** for problem solving

(Abstract) state = set of real states

(Abstract) action = complex combination of real actions

e.g., “Arad → Zerind” represents a complex set  
of possible routes, detours, rest stops, etc.


For guaranteed realizability, **any** real state “in Arad”  
must get to **some** real state “in Zerind”

(Abstract) solution =

set of real paths that are solutions in the real world

Each abstract action should be “easier” than the original problem!

## Example: vacuum world state space graph




states??:

actions??:

goal test??:

path cost??:

## Example: vacuum world state space graph




states??: integer dirt and robot locations (ignore dirt **amounts** etc.)

actions??:

goal test??:

path cost??:

## Example: vacuum world state space graph




states??: integer dirt and robot locations (ignore dirt **amounts** etc.)

actions??: *Left, Right, Suck, NoOp*

goal test??:

path cost??:

## Example: vacuum world state space graph




states??: integer dirt and robot locations (ignore dirt **amounts** etc.)

actions??: *Left, Right, Suck, NoOp*

goal test??: no dirt

path cost??:

## Example: vacuum world state space graph



states??: integer dirt and robot locations (ignore dirt **amounts** etc.)

actions??: *Left*, *Right*, *Suck*, *NoOp*

goal test??: no dirt

path cost??: 1 per action (0 for *NoOp*)



## Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

states??:

actions??:

goal test??:

path cost??:

## Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

states??: integer locations of tiles (ignore intermediate positions)

actions??:

goal test??:

path cost??:

## Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

states??: integer locations of tiles (ignore intermediate positions)

actions??: move blank left, right, up, down (ignore unjamming etc.)

goal test??:

path cost??:

## Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

states??: integer locations of tiles (ignore intermediate positions)

actions??: move blank left, right, up, down (ignore unjamming etc.)

goal test??: = goal state (given)

path cost??:

## Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

states??: integer locations of tiles (ignore intermediate positions)


actions??: move blank left, right, up, down (ignore unjamming etc.)

goal test??: = goal state (given)

path cost??: 1 per move

[Note: optimal solution of  $n$ -Puzzle family is NP-hard]

## Example: robotic assembly



states??: real-valued coordinates of robot joint angles  
parts of the object to be assembled

actions??: continuous motions of robot joints

goal test??: complete assembly **with no robot included!**

path cost??: time to execute


# Tree search algorithms

Basic idea:

offline, simulated exploration of state space  
by generating successors of already-explored states  
(a.k.a. **expanding** states)


```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

# Tree search example






# Tree search example



# Tree search example

Arad

Sibiu Timisoara Zerind →




## Implementation: states vs. nodes

A **state** is a (representation of) a physical configuration

A **node** is a data structure constituting part of a search tree

includes **parent**, **children**, **depth**, **path cost**  $g(x)$

States do not have parents, children, depth, or path cost!



The **EXPAND** function creates new nodes, filling in the various fields and using the **SUCCESSORFN** of the problem to create the corresponding states.

# Implementation: general tree search

front + FIFO end  
↙ ○ ↘

**function** TREE-SEARCH(*problem*, *fringe*) **returns** a solution, or failure

*fringe*  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[*problem*]), *fringe*)

**loop do**

**if** *fringe* is empty **then return** failure

*node*  $\leftarrow$  REMOVE-FRONT(*fringe*)

**if** GOAL-TEST(*problem*, STATE(*node*)) **then return** *node*

*fringe*  $\leftarrow$  INSERTALL(EXPAND(*node*, *problem*), *fringe*)

---

**function** EXPAND(*node*, *problem*) **returns** a set of nodes

*successors*  $\leftarrow$  the empty set

**for each** *action*, *result* **in** SUCCESSOR-FN(*problem*, STATE[*node*]) **do**

*s*  $\leftarrow$  a new NODE

PARENT-NODE[*s*]  $\leftarrow$  *node*; ACTION[*s*]  $\leftarrow$  *action*; STATE[*s*]  $\leftarrow$  *result*

PATH-COST[*s*]  $\leftarrow$  PATH-COST[*node*] + STEP-COST(STATE[*node*], *action*, *result*)

DEPTH[*s*]  $\leftarrow$  DEPTH[*node*] + 1

add *s* to *successors*

**return** *successors*

# Search strategies

A strategy is defined by picking the **order of node expansion**

Strategies are evaluated along the following dimensions:

**completeness**—does it always find a solution if one exists?

**time complexity**—number of nodes generated/expanded

**space complexity**—maximum number of nodes in memory

**optimality**—does it always find a least-cost solution?

Time and space complexity are measured in terms of

$b$ —maximum branching factor of the search tree

$d$ —depth of the least-cost solution

$m$ —maximum depth of the state space (may be  $\infty$ )

# Uninformed search strategies

**Uninformed** strategies use only the information available in the problem definition

Breadth-first search (BFS)

Depth-first search (DFS)

Depth-limited search

Iterative deepening search

Uniform-cost search

# Breadth-first search

**function** BREADTH-FIRST-SEARCH(*problem*) **returns** a solution node or *failure*

*node*  $\leftarrow$  NODE(*problem*.INITIAL)

**if** *problem*.IS-GOAL(*node*.STATE) **then return** *node*

*frontier*  $\leftarrow$  a FIFO queue, with *node* as an element

*reached*  $\leftarrow$  {*problem*.INITIAL}

**while not** IS-EMPTY(*frontier*) **do**

*node*  $\leftarrow$  POP(*frontier*)

**for each** *child* **in** EXPAND(*problem*, *node*) **do**

*s*  $\leftarrow$  *child*.STATE


**if** *problem*.IS-GOAL(*s*) **then return** *child*

**if** *s* is not in *reached* **then**

add *s* to *reached*

add *child* to *frontier*

**return** *failure*




# Breadth-first search

Expand shallowest unexpanded node

**Implementation:**

*fringe* is a FIFO queue, i.e., new successors go at end






# Breadth-first search

Expand shallowest unexpanded node

**Implementation:**

*fringe* is a FIFO queue, i.e., new successors go at end




# Breadth-first search

Expand shallowest unexpanded node

**Implementation:**

*fringe* is a FIFO queue, i.e., new successors go at end




## Breadth-first search

Expand shallowest unexpanded node

**Implementation:**

*fringe* is a FIFO queue, i.e., new successors go at end



Fringe  
A  
BC  
CDE  
DEFG  
EFG  
FG  
G  
∅

# Properties of breadth-first search

Complete?? Yes (if  $b$  is finite)

# Properties of breadth-first search

Complete?? Yes (if  $b$  is finite)

Time??  $b + b^2 + b^3 + \dots + b^d = O(b^d)$ , i.e., exponential in  $d$

# Properties of breadth-first search

Complete?? Yes (if  $b$  is finite)

Time??  $b + b^2 + b^3 + \dots + b^d = O(b^d)$ , i.e., exponential in  $d$

Space??  $O(b^d)$  (keeps every node in memory)

# Properties of breadth-first search

Complete?? Yes (if  $b$  is finite)

Time??  $b + b^2 + b^3 + \dots + b^d = O(b^d)$ , i.e., exponential in  $d$

Space??  $O(b^d)$  (keeps every node in memory)

Optimal?? Yes (if cost = 1 per step); not optimal in general

# Properties of breadth-first search

Complete?? Yes (if  $b$  is finite)

Time??  $b + b^2 + b^3 + \dots + b^d = O(b^d)$ , i.e., exponential in  $d$

Space??  $O(b^d)$  (keeps every node in memory)

Optimal?? Yes (if cost = 1 per step); not optimal in general

**Space** is the big problem; can easily generate nodes at 100MB/sec  
so 24hrs = 8640GB.




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*fringe* = LIFO queue, i.e., put successors at front






# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*fringe* = LIFO queue, i.e., put successors at front



# Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops  
Modify to avoid repeated states along path  
⇒ complete in finite spaces

# Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path

⇒ complete in finite spaces

Time??  $O(b^m)$ : terrible if  $m$  is much larger than  $d$ , where  $m$  is the maximum depth of any node

but if solutions are dense, may be much faster than breadth-first

# Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path

⇒ complete in finite spaces

Time??  $O(b^m)$ : terrible if  $m$  is much larger than  $d$ , where  $m$  is the maximum depth of any node

but if solutions are dense, may be much faster than breadth-first

Space??  $O(bm)$ , i.e., linear space!

# Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path

⇒ complete in finite spaces

Time??  $O(b^m)$ : terrible if  $m$  is much larger than  $d$ , where  $m$  is the maximum depth of any node

but if solutions are dense, may be much faster than breadth-first

Space??  $O(bm)$ , i.e., linear space!

Optimal?? No



# Depth-limited search

= depth-first search with depth limit  $\ell$ ,  
i.e., nodes at depth  $\ell$  have no successors

## Recursive implementation:

**function** ITERATIVE-DEEPENING-SEARCH(*problem*) **returns** a solution node or *failure*  
  **for** *depth* = 0 **to**  $\infty$  **do**  
    *result*  $\leftarrow$  DEPTH-LIMITED-SEARCH(*problem*, *depth*)  
    **if** *result*  $\neq$  *cutoff* **then return** *result*

**function** DEPTH-LIMITED-SEARCH(*problem*,  $\ell$ ) **returns** a node or *failure* or *cutoff*  
  *frontier*  $\leftarrow$  a LIFO queue (stack) with NODE(*problem*.INITIAL) as an element  
  *result*  $\leftarrow$  *failure*  
  **while not** IS-EMPTY(*frontier*) **do**  
    *node*  $\leftarrow$  POP(*frontier*)  
    **if** *problem*.IS-GOAL(*node*.STATE) **then return** *node*  
    **if** DEPTH(*node*) >  $\ell$  **then**  
      *result*  $\leftarrow$  *cutoff*  
    **else if not** IS-CYCLE(*node*) **do**  
      **for each** *child* **in** EXPAND(*problem*, *node*) **do**  
        add *child* to *frontier*  
  **return** *result*





# Iterative deepening search

Limit = 0











# Iterative deepening search

Limit = 1




# Iterative deepening search

Limit = 2



# Iterative deepening search

Limit = 3



# Properties of iterative deepening search

Complete?? Yes

# Properties of iterative deepening search

Complete?? Yes

Time??  $db^1 + (d-1)b^2 + \dots + b^d = O(b^d)$

# Properties of iterative deepening search

Complete?? Yes

Time??  $db^1 + (d-1)b^2 + \dots + b^d = O(b^d)$

Space??  $O(bd)$



# Properties of iterative deepening search

Complete?? Yes

Time??  $db^1 + (d-1)b^2 + \dots + b^d = O(b^d)$

Space??  $O(bd)$

Optimal?? Yes, if step cost = 1

There is some extra cost for generating the upper levels multiple times, but it is not large. E.g., numerical comparison for  $b = 10$  and  $d = 5$ , solution at far right leaf:



$$N(\text{IDS}) = 50 + 400 + 3,000 + 20,000 + 100,000 = 123,450$$

$$\begin{aligned} N(\text{BFS}) &= 10 + 100 + 1,000 + 10,000 + 100,000 \\ &= 111,110 \end{aligned}$$

IDS does better because other nodes at depth  $d$  are not expanded  
BFS can be modified to apply goal test when a node is **generated**

## Repeated states

Failure to detect repeated states can turn a linear problem into an exponential one!



# Graph search

```
function GRAPH-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  initialize the explored set to be empty
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    add the node to the explored set
    expand the chosen node, adding the resulting nodes to the frontier
    only if not in the frontier or explored set
  end
```

# Uniform-cost search

- ▶ When all step costs are equal, breadth-first search is optimal

# Uniform-cost search

- ▶ When all step costs are equal, breadth-first search is optimal
- ▶ What if all step costs are not equal?

# Uniform-cost search

**function** UNIFORM-COST-SEARCH(*problem*) **returns** a solution, or failure

*node*  $\leftarrow$  a node with STATE = *problem*.INITIAL-STATE, PATH-COST = 0

*frontier*  $\leftarrow$  a priority queue ordered by PATH-COST, with *node* as the only element

*explored*  $\leftarrow$  an empty set

**loop do**

**if** EMPTY?(*frontier*) **then return** failure

*node*  $\leftarrow$  POP(*frontier*) /\* chooses the lowest-cost node in *frontier* \*/

**if** *problem*.GOAL-TEST(*node*.STATE) **then return** SOLUTION(*node*)

    add *node*.STATE to *explored*

**for each** *action* **in** *problem*.ACTIONS(*node*.STATE) **do**

*child*  $\leftarrow$  CHILD-NODE(*problem*, *node*, *action*)


**if** *child*.STATE is not in *explored* or *frontier* **then**

*frontier*  $\leftarrow$  INSERT(*child*, *frontier*)


**else if** *child*.STATE is in *frontier* with higher PATH-COST **then**

            replace that *frontier* node with *child*

## Uniform-cost search



## Uniform-cost search



Sibiu → Rimnicu Vilcea → Fagaras → Pitesti → Bucharest



# Uniform-cost search

Expand least-cost unexpanded node

Let  $g(n)$  be the sum of the cost (path cost) from start to node  $n$

**Implementation:**

*fringe* = queue ordered by path cost, lowest first

Equivalent to breadth-first if step costs all equal

Complete?? Yes, if step cost  $\geq \epsilon$

Time?? # of nodes with  $g \leq$  cost of optimal solution,  $O(b^{\lceil C^*/\epsilon \rceil})$   
where  $C^*$  is the cost of the optimal solution

Space?? # of nodes with  $g \leq$  cost of optimal solution,  $O(b^{\lceil C^*/\epsilon \rceil})$

Optimal?? Yes—nodes expanded in increasing order of  $g(n)$

## Summary of algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes*	Yes*	No	Yes, if $l \geq d$	Yes
Time	$b^d$	$b^{\lceil C^*/\epsilon \rceil}$	$b^m$	$b^l$	$b^d$
Space	$b^d$	$b^{\lceil C^*/\epsilon \rceil}$	$bm$	$bl$	$bd$
Optimal?	Yes*	Yes	No	No	Yes*

# Summary

- Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored
- Variety of uninformed search strategies
- Iterative deepening search uses only linear space and not much more time than other uninformed algorithms
- Graph search can be exponentially more efficient than tree search

# Acknowledgment

The slides are adapted from Stuart Russell et al.